

# Automatic Transformations for Communication-Minimized Parallelization and Locality Optimization in the Polyhedral Model

Uday Bondhugula

Department of Computer Science & Engineering  
The Ohio State University

U. Bondhugula, M. Baskaran, S. Krishnamoorthy, A. Rountev, P. Sadayappan  
(OSU), J. Ramanujam (LSU)

April 3, 2008

- 1 Introduction
- 2 Polyhedral techniques for program optimization
- 3 A new transformation framework
- 4 Implementation
- 5 Related and Future work

# Multicore architectures

- Architectures with multiple processing units on chip have become mainstream
  - General-purpose multicore microprocessors
  - Specialized: GPUs, Cell, MPSoCs
- Difficulty of Parallel Programming
- **Automatic Parallelization:** user does nothing

# Multicore architectures


- Architectures with multiple processing units on chip have become mainstream
  - General-purpose multicore microprocessors
  - Specialized: GPUs, Cell, MPSoCs
- Difficulty of Parallel Programming
- **Automatic Parallelization:** user does nothing

# Multicore architectures


- Architectures with multiple processing units on chip have become mainstream
  - General-purpose multicore microprocessors
  - Specialized: GPUs, Cell, MPSoCs
- Difficulty of Parallel Programming
- **Automatic Parallelization:** user does nothing

- 1 Introduction
- 2 Polyhedral techniques for program optimization
- 3 A new transformation framework
- 4 Implementation
- 5 Related and Future work

# Background: polyhedral/polytope model




(a) A hyperplane




(b) A polytope

- Loop nests with regular accesses (statically predictable) - sequences of imperfectly nested loops
- More general code like non-affine accesses, dynamic control can also be handled with conservative assumptions

# Background: polyhedral/polytope model



(c) A hyperplane



(d) A polytope

- Loop nests with regular accesses (statically predictable) - sequences of imperfectly nested loops
- More general code like non-affine accesses, dynamic control can also be handled with conservative assumptions



# Polyhedral compiler framework

- ① Dependence analysis (exact affine dependences) [Feautrier91, Pugh92, Vasilache06ICS]
  - ② Automatic transformations [Feautrier92, Lim/Lam97, Griebel04]
  - ③ Code generation from specified transforms [Omega90s, Quilleré00, **Bastoul04**, **CLooG**, Vasilache06CC]
- Significant advances in first and last step during this decade
  - Semi-automatic approaches demonstrating polyhedral model as a powerful representation [Cohen05ICS, Girbal06IJPP]

# Polyhedral compiler framework

- ① Dependence analysis (exact affine dependences) [Feautrier91, Pugh92, Vasilache06ICS]
  - ② Automatic transformations [Feautrier92, Lim/Lam97, Griebel04]
  - ③ Code generation from specified transforms [Omega90s, Quilleré00, **Bastoul04**, **CLooG**, Vasilache06CC]
- Significant advances in first and last step during this decade
  - Semi-automatic approaches demonstrating polyhedral model as a powerful representation [Cohen05ICS, Girbal06IJPP]

# Polyhedral optimization

- ① Dependence analysis (exact affine dependences) [Feautrier91, Pugh92, Vasilache06ICS]
- ② **Automatic transformations (for parallelism and locality)**
- ③ Code generation from specified transforms [Omega90s, Quilleré00, Bastoul04, CLooG, Vasilache06CC]

Our work: a new theoretical framework for automatic transformation

# Polyhedral model: an example

```

for (i=0; i<N; i++)
  for (j=0; j<N; j++)
    S1: A[i,j] = A[i,j]+u1[i]*v1[j] + u2[i]*v2[j];

for (i''=0; i''<N; i''++)
  for (j'=0; j'<N; j'++)
    S2: x[i''] = x[i''] + A[j', i''] * y[j'];

```



(2) The Generalized Dependence Graph

$$D^{S_1} \left( \begin{pmatrix} 1 & 0 & 0 & 0 \\ 0 & 1 & 0 & 0 \\ -1 & 0 & 1 & -1 \\ 0 & -1 & 1 & -1 \end{pmatrix} \right) \begin{pmatrix} i \\ j \\ N \\ 1 \end{pmatrix} \geq 0$$

(1.1) Statement domain

$$\left( \begin{array}{cccccc|c} 1 & 0 & 0 & 0 & 0 & 0 & i \\ 0 & -1 & 0 & 0 & 1 & -1 & j \\ 0 & 0 & 1 & 0 & 0 & 0 & i' \\ 0 & 0 & 0 & -1 & 1 & -1 & j' \\ \hline 1 & 0 & 0 & -1 & 0 & 0 & N \\ 0 & 1 & -1 & 0 & 0 & 0 & 1 \end{array} \right) \begin{matrix} \geq 0 \\ \geq 0 \\ \geq 0 \\ \geq 0 \\ = 0 \\ = 0 \end{matrix}$$

(3.1) An exact dependence polyhedron ( $S_1 \rightarrow S_2$ )

# Polyhedral model: Motivation for automatic transformation

## GEMVER

```

dcopy(m * n, A, B, 1);
dger(m, n, 1.0, u1, 1, v1, 1, B, m);
dger(m, n, 1.0, u2, 1, v2, 1, B, m);
dcopy(n, z, x, 1);
dgemv('T', m, n, beta, B, m, y, 1, 1.0, x, 1);
dgemv('N', m, n, alpha, B, m, x, 1, 0.0, w, 1);

```

BLAS version [Siek et al., POHLL '08]

$$\begin{aligned}
 B &= A + u_1 v_1^T + u_2 v_2^T \\
 x &= \beta B^T y + z \\
 w &= \alpha Bx
 \end{aligned}$$



The Generalized Dependence Graph

$$\begin{pmatrix}
 1 & 0 & 0 & 0 & 0 & 0 \\
 0 & -1 & 0 & 0 & 1 & -1 \\
 0 & 0 & 1 & 0 & 0 & 0 \\
 0 & 0 & 0 & -1 & 1 & -1 \\
 \hline
 1 & 0 & 0 & -1 & 0 & 0 \\
 0 & 1 & -1 & 0 & 0 & 0
 \end{pmatrix}
 \begin{bmatrix}
 i \\
 j \\
 i' \\
 j' \\
 N \\
 1
 \end{bmatrix}
 \begin{array}{l}
 \geq 0 \\
 \geq 0 \\
 \geq 0 \\
 \geq 0 \\
 = 0 \\
 = 0
 \end{array}$$

An exact dependence polyhedron ( $S1 \rightarrow S2$ )

# Polyhedral model: Motivation for automatic transformation

## GEMVER

```

for (i=0; i<N; i++)
  for (j=0; j<N; j++)
    S1: A[i,j] = A[i,j]+u1[i]*v1[j] + u2[i]*v2[j];

for (i'=0; i'<N; i'++)
  for (j'=0; j'<N; j'++)
    S2: x[i'] = x[i']+A[j',i']*y[j'];
    original code
  
```



The Generalized Dependence Graph

	S1			S2			
	<i>i</i>	<i>j</i>	<i>const</i>	<i>i</i>	<i>j</i>	<i>const</i>	
<i>c</i> <sub>1</sub>	0	1	0	1	0	0	parallel
<i>c</i> <sub>2</sub>	1	0	0	0	1	0	fwd_dep
<i>c</i> <sub>3</sub>	0	0	0	0	0	1	scalar

Statement-wise transformation

```

for (c1=0; c1<N; c1++)
  for (c2=0; c2<N; c2++)
    A[c2,c1] = A[c2,c1]+u[c2]*v[c1];
    x[c1] = x[c1]+A[c2,c1]*y[c2];
  
```

Transformed code (not final)

$$\begin{pmatrix} 1 & 0 & 0 & 0 & 0 & 0 \\ 0 & -1 & 0 & 0 & 1 & -1 \\ 0 & 0 & 1 & 0 & 0 & 0 \\ 0 & 0 & 0 & -1 & 1 & -1 \\ \hline 1 & 0 & 0 & -1 & 0 & 0 \\ 0 & 1 & -1 & 0 & 0 & 0 \end{pmatrix} \begin{bmatrix} i \\ j \\ i' \\ j' \\ N \\ 1 \end{bmatrix} \begin{matrix} \geq 0 \\ \geq 0 \\ \geq 0 \\ \geq 0 \\ = 0 \\ = 0 \end{matrix}$$

An exact dependence polyhedron ( $S1 \rightarrow S2$ )

Cores	Our (poly)	native cc	ACML 4.0.1/ifort
1	0.348s	2.33s	0.679s
2	0.238s	1.46s	0.59s

AMD Opteron (dual core) 2.6 GHz, execution time

# Polyhedral model: Motivation for automatic transformation

## GEMVER

```

for (i=0; i<N; i++)
  for (j=0; j<N; j++)
    S1: A[i,j] = A[i,j]+u1[i]*v1[j] + u2[i]*v2[j];

for (i'=0; i'<N; i'++)
  for (j'=0; j'<N; j'++)
    S2: x[i'] = x[i']+A[j',i']*y[j'];
    original code
  
```



The Generalized Dependence Graph

	S1			S2			
	i	j	const	i	j	const	
c <sub>1</sub>	0	1	0	1	0	0	parallel
c <sub>2</sub>	1	0	0	0	1	0	fwd_dep
c <sub>3</sub>	0	0	0	0	0	1	scalar

Statement-wise transformation

$$\begin{pmatrix}
 1 & 0 & 0 & 0 & 0 & 0 \\
 0 & -1 & 0 & 0 & 1 & -1 \\
 0 & 0 & 1 & 0 & 0 & 0 \\
 0 & 0 & 0 & -1 & 1 & -1 \\
 1 & 0 & 0 & -1 & 0 & 0 \\
 0 & 1 & -1 & 0 & 0 & 0
 \end{pmatrix}
 \begin{bmatrix}
 i \\
 j \\
 j' \\
 i' \\
 N \\
 1
 \end{bmatrix}
 \begin{matrix}
 \geq 0 \\
 \geq 0 \\
 \geq 0 \\
 \geq 0 \\
 = 0 \\
 = 0
 \end{matrix}$$

An exact dependence polyhedron (S1→S2)

```

for (c1=0; c1<N; c1++)
  for (c2=0; c2<N; c2++)
    A[c2,c1] = A[c2,c1]+u[c2]*v[c1];
    x[c1] = x[c1]+A[c2,c1]*y[c2];
  
```

Transformed code (not final)

Cores	Over native cc	Over vendor BLAS
1	6.7x	2.0x
2	6.1x	2.5x

AMD Opteron (dual core) 2.6 GHz, Poly Speedup

# Affine Transformations in the polyhedral model

A one-dimensional affine transform for statement  $S_k$  is defined by:

$$\begin{aligned}\phi_{S_k}(\vec{i}) &= \begin{bmatrix} c_1 & c_2 & \dots & c_{m_{S_k}} \end{bmatrix} \begin{pmatrix} \vec{i} \end{pmatrix} + c_0 \\ &= \begin{bmatrix} c_1 & c_2 & \dots & c_{m_{S_k}} & c_0 \end{bmatrix} \begin{pmatrix} \vec{i} \\ 1 \end{pmatrix}\end{aligned}$$

where  $[c_0, c_1, c_2, \dots, c_{m_{S_k}}] \in \mathbb{Z}$ .

- An affine transform  $\equiv$  A new scanning hyperplane  $\equiv$  A loop in the transformed space (with a particular property)



# Affine transformations

- A transformation for each statement,  $S$ :  $T_S \vec{i} + \vec{b}_S$

$$\begin{pmatrix} i'_1 \\ i'_2 \\ i'_3 \\ \vdots \\ i'_n \end{pmatrix} = \begin{pmatrix} c_{11} & c_{12} & \dots & c_{1n} \\ c_{21} & c_{22} & \dots & c_{2n} \\ \vdots & \vdots & \vdots & \vdots \\ c_{n1} & c_{n2} & \dots & c_{nn} \end{pmatrix} \begin{pmatrix} i_1 \\ i_2 \\ i_3 \\ \vdots \\ i_n \end{pmatrix} + \begin{pmatrix} c_{01} \\ c_{02} \\ c_{03} \\ \vdots \\ c_{0n} \end{pmatrix}$$

- Full column-ranked transform is a one-to-one mapping
- Each 1-d transform ( $\phi$ ) can later be marked as a space loop or time (sequential) loop or a band of them can be tiled
- **Problem:** how do you find good transformations optimized for parallelism and locality?

# Tiling in the polyhedral model

- Tile: a portion of the iteration space that can be executed atomically
- **Tiling for parallelism:** Enables coarse-grained parallelization: reduces frequency of communication
- **Tiling for locality:** Allows reuse along multiple dimensions - tile fits in faster memory
- Tile shape and size affect the volume and frequency of communication / number of cache misses
- Legality of tiling for restricted input and/or weaker dependence abstractions are well understood [Irigoin and Triolet 88, Wolf/Lam 91, Darte et al. 97]

# Tiling in the polyhedral model

- Tile: a portion of the iteration space that can be executed atomically
- **Tiling for parallelism:** Enables coarse-grained parallelization: reduces frequency of communication
- **Tiling for locality:** Allows reuse along multiple dimensions - tile fits in faster memory
- Tile shape and size affect the volume and frequency of communication / number of cache misses
- Legality of tiling for restricted input and/or weaker dependence abstractions are well understood [Irigoin and Triolet 88, Wolf/Lam 91, Darte et al. 97]

# Tiling in the polyhedral model

- Tile: a portion of the iteration space that can be executed atomically
- **Tiling for parallelism:** Enables coarse-grained parallelization: reduces frequency of communication
- **Tiling for locality:** Allows reuse along multiple dimensions - tile fits in faster memory
- Tile shape and size affect the volume and frequency of communication / number of cache misses
- Legality of tiling for restricted input and/or weaker dependence abstractions are well understood [Irigoin and Triolet 88, Wolf/Lam 91, Darte et al. 97]

# Tiling in the polyhedral model

- Tile: a portion of the iteration space that can be executed atomically
- **Tiling for parallelism:** Enables coarse-grained parallelization: reduces frequency of communication
- **Tiling for locality:** Allows reuse along multiple dimensions - tile fits in faster memory
- Tile shape and size affect the volume and frequency of communication / number of cache misses
- Legality of tiling for restricted input and/or weaker dependence abstractions are well understood [Irigoin and Triolet 88, Wolf/Lam 91, Darte et al. 97]

# Tiling and its legality for exact polyhedral dependences

If  $\vec{s}$  and  $\vec{t}$  are dependent through dependence polyhedron  $P_e$  (corresponding to a dependence edge  $s_i \rightarrow s_j$ ), then

$$\phi_{s_i}(\vec{t}) - \phi_{s_j}(\vec{s}) \geq 0, \quad \langle \vec{s}, \vec{t} \rangle \in P_e$$

- Extension of classic condition from Irigoin and Triolet [PoPL88]: dependence only has non-negative components along  $\phi$
- At least two independent  $\phi$ 's that satisfy the above property for all unsatisfied dependences so far  $\rightarrow$  Tiling
- For affine dependences and statements of different dimensionalities (coming from arbitrarily nested loops)

# Tiling and its legality for exact polyhedral dependences

If  $\vec{s}$  and  $\vec{t}$  are dependent through dependence polyhedron  $P_e$  (corresponding to a dependence edge  $s_i \rightarrow s_j$ ), then

$$\phi_{s_i}(\vec{t}) - \phi_{s_j}(\vec{s}) \geq 0, \quad \langle \vec{s}, \vec{t} \rangle \in P_e$$

- Extension of classic condition from Irigoin and Triolet [PoPL88]: dependence only has non-negative components along  $\phi$
- At least two independent  $\phi$ 's that satisfy the above property for all unsatisfied dependences so far  $\rightarrow$  Tiling
- For affine dependences and statements of different dimensionalities (coming from arbitrarily nested loops)

# Capturing communication volume and reuse distance

- $\phi_{S_i}(\vec{t}) - \phi_{S_j}(\vec{s})$  is a very important affine function
- Define an affine form  $\delta_e$  in  $\vec{s}, \vec{t}$  (for every dependence):

$$\delta_e(\vec{s}, \vec{t}) = \phi_{S_i}(\vec{t}) - \phi_{S_j}(\vec{s}), \quad \langle \vec{s}, \vec{t} \rangle \in P_e$$

- Dot product of hyperplane with dependence ( $\mathbf{h} \cdot \vec{d}$ )
- Number of hyperplane instances separating source and sink



# Capturing communication volume and reuse distance


- $\phi_{S_i}(\vec{t}) - \phi_{S_j}(\vec{s})$  is a very important affine function
- Define an affine form  $\delta_e$  in  $\vec{s}, \vec{t}$  (for every dependence):

$$\delta_e(\vec{s}, \vec{t}) = \phi_{S_i}(\vec{t}) - \phi_{S_j}(\vec{s}), \quad \langle \vec{s}, \vec{t} \rangle \in P_e$$

- Dot product of hyperplane with dependence ( $\mathbf{h} \cdot \vec{d}$ )
- Number of hyperplane instances separating source and sink

# Capturing communication volume and reuse distance

- Consider the stencil code below with deps:  $(1,0)$ ,  $(1,1)$ ,  $(1,-1)$



- Represents the component of a dependence along the hyperplane  $(\phi)$ 
  - Communication volume (per unit area) at processor tile boundaries
  - Cache misses at local tile edges (L2, L1, registers)

# Cost function in the polyhedral framework

Minimizing  $\delta_e(\vec{s}, \vec{t})$  can be used to:

- Find hyperplane that minimizes inter-tile communication volume (rate per unit area)
- Find direction that minimizes reuse distance

But directly attempting to optimize  $\delta_e$  is problematic

- Not expressible as a linear function of transformation coefficients
- $\phi(\vec{t}) - \phi(\vec{s})$  could be  $c_1 i + (c_2 - c_3)j$ , where  $1 \leq i \leq N \wedge 1 \leq j \leq N \wedge i \leq j$

# Using a bounding function

- Even if  $\delta_e$  involves loop variables, it will still be less than a linear function of parameters

$$\begin{aligned}\phi_{s_i}(\vec{t}) - \phi_{s_j}(\vec{s}) &\leq \mathbf{u} \cdot \vec{n} + w, \quad \langle \vec{s}, \vec{t} \rangle \in P_e \\ v(\vec{n}) - \delta_e(\vec{s}, \vec{t}) &\geq 0, \quad \langle \vec{s}, \vec{t} \rangle \in P_e, \quad \forall e \in E\end{aligned}$$

where  $\vec{n}$  is the vector of program parameters

- Bound from above and minimize the coefficients of the bound after linearizing with Affine form of the Farkas Lemma
- The  $\delta_e$  for each dependence is bounded in this manner

# Using a bounding function

- Even if  $\delta_e$  involves loop variables, it will still be less than a linear function of parameters

$$\begin{aligned} \phi_{s_i}(\vec{t}) - \phi_{s_j}(\vec{s}) &\leq \mathbf{u} \cdot \vec{n} + w, \quad \langle \vec{s}, \vec{t} \rangle \in P_e \\ v(\vec{n}) - \delta_e(\vec{s}, \vec{t}) &\geq 0, \quad \langle \vec{s}, \vec{t} \rangle \in P_e, \quad \forall e \in E \end{aligned}$$

where  $\vec{n}$  is the vector of program parameters

- Bound from above and minimize the coefficients of the bound after linearizing with Affine form of the Farkas Lemma
- The  $\delta_e$  for each dependence is bounded in this manner

# Bounding function approach (Farkas Lemma)

Now, use the affine form of the Farkas lemma on the bounding constraint

$$\mathbf{u} \cdot \vec{n} + w - \delta_e(\vec{s}, \vec{t}) \equiv \lambda_{e0} + \sum_{k=1}^{m_e} \lambda_{ek} * f_{ek}(\vec{s}, \vec{t}), \quad f_{ek} \in P_e$$

- Linearizes legality and communication volume/reuse distance bounding constraints
- Use PIP to find the lexicographic minimal solution that satisfies above system with  $\vec{u}$  and  $w$  in the leading position

$$\text{minimize}_{\prec} \{ \mathbf{u}, w, \dots, c'_i s, \dots \}$$

- Minimizes the maximum dependence component along hyperplane normal (across all dependences)

# Implications of cost function minimization

- $u = 0, w = 0$ :  $\phi$  is a communication-free parallel hyperplane
  - $u = 0, w = \text{const}$ : constant minimum boundary line  
communication/cache misses
  - $u > 0$ : non-constant (large) amount of communication/cache misses
- 
- A solution for multiple statements at a level is a fused loop
  - Several refinements possible

# Implications of cost function minimization

- $u = 0, w = 0$ :  $\phi$  is a communication-free parallel hyperplane
- $u = 0, w = \text{const}$ : constant minimum boundary line  
communication/cache misses
- $u > 0$ : non-constant (large) amount of communication/cache misses
- A solution for multiple statements at a level is a fused loop
- Several refinements possible



# Implications of cost function minimization

- $u = 0, w = 0$ :  $\phi$  is a communication-free parallel hyperplane
- $u = 0, w = \text{const}$ : constant minimum boundary line  
communication/cache misses
- $u > 0$ : non-constant (large) amount of communication/cache misses
- A solution for multiple statements at a level is a fused loop
- Several refinements possible

# Finding independent solutions iteratively

- Once a solution (a hyperplane for all statements is found), the same formulation is augmented with additional linear independence constraints
- Find independent solutions one after the other till enough hyperplanes to scan all statements are found
- Dependences are not removed as we proceed from hyperplane to another unless they need to be
  - A hierarchy of permutable loop nest sets is found (with a cost function and for any polyhedral input)
  - Fully tilable, partially tilable (tiling at any arbitrary level) or inner tilable loops are identified

# Finding independent solutions iteratively

- Once a solution (a hyperplane for all statements is found), the same formulation is augmented with additional linear independence constraints
- Find independent solutions one after the other till enough hyperplanes to scan all statements are found
- Dependences are not removed as we proceed from hyperplane to another unless they need to be
  - A hierarchy of permutable loop nest sets is found (with a cost function and for any polyhedral input)
  - Fully tilable, partially tilable (tiling at any arbitrary level) or inner tilable loops are identified

# Summary of Algorithm

- Affine dependences are pushed as much inside as possible
- Outer loops are space with minimal communication
  - Synchronization-free parallel loops end up first if they exist
  - Next, space loops with minimal communication are found
- Inner loops are sequential time with maximum reuse
- Linearly independent sub-space construction, avoiding trivial solutions: reasonable choices made to avoid combinatorial explosion
- **Complexity:** Very fast in practice

# Summary of Algorithm

- Affine dependences are pushed as much inside as possible
- Outer loops are space with minimal communication
  - Synchronization-free parallel loops end up first if they exist
  - Next, space loops with minimal communication are found
- Inner loops are sequential time with maximum reuse
- Linearly independent sub-space construction, avoiding trivial solutions: reasonable choices made to avoid combinatorial explosion
- **Complexity:** Very fast in practice

- 1 Introduction
- 2 Polyhedral techniques for program optimization
- 3 A new transformation framework
- 4 Implementation**
- 5 Related and Future work

# The PLuTo system

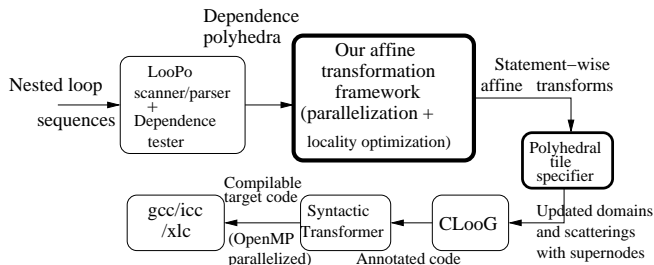


Figure: The PLuTo automatic parallelizer (source-to-source)

- Framework implemented with PipLib and interfaced with LooPo frontend (Univ. of Passau, Germany) and CLOoG (Cédric Bastoul, INRIA Saclay)

# Transformation framework running time

Code	Num of statements	Num of loops	Num of deps	Running time
2-d Jacobi	2	6	20	0.05s
Haar 1-d	3	5	12	0.018s
LU	2	5	10	0.022s
TCE 4-index	4	20	15	0.20s
Swim	58	160	639	20.9s

Table: Transformation tool running time



# Experimental results: preview

- Intel Core2 Quad Q6600 2.4 GHz (quad core), DDR2 667 RAM
- 32 KB L1 cache, 8 MB (shared) L2 cache (4MB per core pair)
- ICC 10.1 (-fast), Linux 2.6.18 x86-64

# Summary of performance improvement

**Table:** Improvement over state-of-the-art research compiler frameworks and native production compiler

Benchmark	Single core improvement		Multi-core speedup (4 cores)	
	over native compiler	over state-of-the-art research	over native compiler	over state-of-the-art research
Jacobi stencil (imperfect)	5.23x	2.1x	20x	2.7x
2-d FDTD	3.7x	3.1x	7.4x	2.5x
3-d Gauss-Seidel	1.6x	1.1x	4.5x	1.5x
LU decomposition	5.6x	5.7x	14x	3.8x
Matrix Vec Transpose	9.3x	5.5x	13x	7x

- High speedups due to simultaneous optimization for parallelism and locality

# Summary of performance improvement

**Table:** Improvement over state-of-the-art research compiler frameworks and native production compiler

Benchmark	Single core improvement		Multi-core speedup (4 cores)	
	over native compiler	over state-of-the-art research	over native compiler	over state-of-the-art research
Jacobi stencil (imperfect)	5.23x	2.1x	20x	2.7x
2-d FDTD	3.7x	3.1x	7.4x	2.5x
3-d Gauss-Seidel	1.6x	1.1x	4.5x	1.5x
LU decomposition	5.6x	5.7x	14x	3.8x
Matrix Vec Transpose	9.3x	5.5x	13x	7x

- High speedups due to simultaneous optimization for parallelism and locality

- 1 Introduction
- 2 Polyhedral techniques for program optimization
- 3 A new transformation framework
- 4 Implementation
- 5 Related and Future work

# Related work

- Previous cost functions developed were for restrictive cases (like single perfectly-nested loops)
- Schedule-based approaches [Feautrier92, Darte/Vivien95, Griebel04 habilitation thesis]
- Affine partitioning [Lim/Lam PoPL'97, ICS'01] - minimizing order of synchronization is not sufficient; Ahmed/Pingali [IJPP'01] (heuristic scalability/practicality)
- Semi-automatic - URUK/WRAP-IT [Cohen05ICS, Girbal06IJPP] - very powerful flexible application of transformations specified manually by an expert

# Related work

- Previous cost functions developed were for restrictive cases (like single perfectly-nested loops)
- Schedule-based approaches [Feautrier92, Darte/Vivien95, Griebel04 habilitation thesis]
- Affine partitioning [Lim/Lam PoPL'97, ICS'01] - minimizing order of synchronization is not sufficient; Ahmed/Pingali [IJPP'01] (heuristic scalability/practicality)
- Semi-automatic - URUK/WRAP-IT [Cohen05ICS, Girbal06IJPP] - very powerful flexible application of transformations specified manually by an expert

# Related work

- Previous cost functions developed were for restrictive cases (like single perfectly-nested loops)
- Schedule-based approaches [Feautrier92, Darte/Vivien95, Griehl04 habilitation thesis]
- Affine partitioning [Lim/Lam PoPL'97, ICS'01] - minimizing order of synchronization is not sufficient; Ahmed/Pingali [IJPP'01] (heuristic scalability/practicality)
- Semi-automatic - URUK/WRAP-IT [Cohen05ICS, Girbal06IJPP] - very powerful flexible application of transformations specified manually by an expert

# Conclusions

- Automatically finding good transforms for imperfectly nested loop sequences for coarse-grained parallelism and locality as is needed in practice
- A beta release of PLuTo (0.0.1) is available  
*<http://pluto-compiler.sourceforge.net>*



# Conclusions

- Automatically finding good transforms for imperfectly nested loop sequences for coarse-grained parallelism and locality as is needed in practice
- A beta release of PLuTo (0.0.1) is available  
*<http://pluto-compiler.sourceforge.net>*

# Future work

- Fusion/parallelization trade-off (which dependences between strongly-connected components to include/cut), interactions of fusion with tiling and prefetching
- Combine with stronger cost models for tile size selection
- Conservative dependence polyhedra for non-affine programs

# Future work

- Fusion/parallelization trade-off (which dependences between strongly-connected components to include/cut), interactions of fusion with tiling and prefetching
- Combine with stronger cost models for tile size selection
- Conservative dependence polyhedra for non-affine programs

# Acknowledgments

- Cédric Bastoul (Université Paris-Sud 11, INRIA Futurs) for **CLooG**, Paul Feautrier and all other contributors of CLooG and PipLib
- Martin Griebl (FMI, Universität Passau, Germany) and team for **LooPo**
- CC'08 reviewers


# References

- ① Automatic Transformations for Communication-Minimized Parallelization and Locality Optimization in the Polyhedral Model  
Uday Bondhugula, M. Baskaran, S. Krishnamoorthy, J. Ramanujam, A. Rountev, and P. Sadayappan.  
OSU CISRC-TR43
- ② A Practical and Fully Automatic Polyhedral Parallelizer and Locality Optimizer. Uday Bondhugula, A. Hartono, J. Ramanujam, and P. Sadayappan.  
ACM SIGPLAN PLDI'08, Jun 2008 (to appear).  
OSU-CISRC-TR70.


# Thank you for your attention

- Questions?

# 2-d Finite Difference Time Domain



(a) Single core:  $T=500$



(b) Parallel:  $n_x = n_y = 2000$ ,  
 $t_{max} = 500$

Figure: 2-d FDTD

# Affine form of the Farkas Lemma

Let the hyperplanes bounding the polytope of statement  $S_k$  be given by:

$$a_{S,k} \left( \begin{array}{c} \vec{i} \\ \vec{n} \end{array} \right) + b_{S,k} \geq 0, \quad k = 1, m_S$$

where  $\vec{n}$  is a vector of the structure parameters. A well-known known result useful in the context of the polytope model is the affine form of the Farkas lemma.

## Lemma (Affine form of Farkas Lemma)

*Let  $\mathcal{D}$  be a non-empty polyhedron defined by  $p$  affine inequalities or faces*

$$a_k \cdot x + b_k \geq 0, \quad k = 1, p$$

*Then, an affine form  $\psi$  is non-negative everywhere in  $\mathcal{D}$  iff it is a positive affine combination of the faces:*

$$\psi(x) \equiv \lambda_0 + \sum_k \lambda_k (a_k x + b_k), \quad \lambda \geq 0$$



# Algorithm

**Input** Generalized dependence graph  $G = (V, E)$  (includes dependence polyhedra  $P_e, e \in E$ )  
 $S_{max}$ : statement with maximum domain dimensionality  
**for** each dependence  $e \in E$  **do**  
    Build legality constraints: apply Farkas Lemma on  $\phi(\vec{t}) - \phi(f_e(\vec{t})) \geq 0$  under  $\vec{t} \in P_e$ , and eliminate all Farkas multipliers  
    Build communication volume/reuse distance bounding constraints: apply Farkas Lemma to  $v(\vec{n}) - (\phi(\vec{t}) - \phi(f(\vec{t}))) \geq 0$  under  $\vec{t} \in P_e$ , and eliminate all Farkas multipliers  
    Aggregate constraints from both into  $C_e(i)$   
**end for**  
**repeat**  
     $C = \emptyset$   
    **for** each dependence edge  $e \in E$  **do**  
         $C \leftarrow C \cup C_e(i)$   
    **end for**  
    Compute lexicographic minimal solution with  $u'$ 's coefficients in the leading position followed by  $w$  to iteratively find independent solutions to  $C$  (orthogonality constraints are added as each soln is found)  
    **if** no solutions were found **then**  
        Cut dependences between two strongly-connected components in the GDG and insert the appropriate *splitter* in the transformation matrices of the statements  
    **end if**  
    Compute  $E_C$ : dependences carried by solutions of Step 12/14; update necessary dependence polyhedra (when a portion of it is satisfied)  
     $E \leftarrow E - E_C$ ; reform the GDG  $(V, E)$   
**until**  $H_{S_{max}}^\perp = \mathbf{0}$  and  $E = \emptyset$   
**Output** A transformation matrix for each statement (with the same number of rows)

# Target architectures

- Transformation framework can be targeted towards general-purpose or special purpose multi-cores (current multicore processors)
- Different kinds/levels, and number of degrees of parallelism required

```

for (t=0; t<tmax; t++) {
  for (j=0; j<ny; j++)
    ey[0][j] = exp(-coeff0*t1);

  for (i=1; i<nx; i++)
    for (j=0; j<ny; j++)
      ey[i][j] = ey[i][j] -
        coeff1*(hz[i][j]-hz[i-1][j]);

  for (i=0; i<nx; i++)
    for (j=1; j<ny; j++)
      ex[i][j] = ex[i][j]
        - coeff1*(hz[i][j]-hz[i][j-1]);

  for (i=0; i<nx; i++)
    for (j=0; j<ny; j++)
      hz[i][j] = hz[i][j] -
        coeff2*(ex[i][j+1]-ex[i][j]
          +ey[i+1][j]-ey[i][j]);
}

```

$$\begin{bmatrix} 1 & 0 & 0 \\ 1 & 1 & 0 \\ 1 & 0 & 0 \end{bmatrix}$$

$$\begin{bmatrix} 1 & 0 & 0 & 0 \\ 1 & 0 & 1 & 0 \\ 1 & 1 & 0 & 0 \end{bmatrix}$$

$$\begin{bmatrix} 1 & 0 & 0 & 0 \\ 1 & 0 & 1 & 0 \\ 1 & 1 & 0 & 0 \end{bmatrix}$$

$$\begin{bmatrix} 1 & 0 & 0 & 0 \\ 1 & 0 & 1 & 1 \\ 1 & 1 & 0 & 1 \end{bmatrix}$$

Figure: 2-d Finite Difference Time Domain code